

Assignment 2 Report

Simple Search Engine using Hadoop MapReduce, Cassandra, and Spark RDD

Course: Big Data - IU

This report presents the implementation of a simple search engine built on top of Hadoop MapReduce, Cassandra, and PySpark RDD. The goal of the assignment was to implement a complete indexing and ranking pipeline for a corpus of plain text documents and to rank documents with BM25 for user queries. The final system supports corpus preparation, distributed index construction, storage of index structures in Cassandra, and query-time ranking with Spark on YARN.

1. Methodology

1.1 Overall System Design

The final system is organized as an end-to-end pipeline with four major stages:

- Local plain text documents are uploaded to HDFS.
- PySpark prepares a normalized one-partition indexing input in HDFS.
- Hadoop Streaming MapReduce jobs build the search index and corpus statistics.
- The index is loaded into Cassandra and queried by a Spark-based BM25 ranker.

The repository uses Docker Compose to start the full environment. The main services are:

- `cluster-master` for Hadoop master services, Spark driver execution, and orchestration.
- `cluster-slave-1` for distributed Hadoop and Spark execution.
- `cassandra-server` for persistent storage of the final index.

The repository contains two corpus modes:

- `app/data_100` is a fixed debug subset of 100 documents. It is used for fast iteration and for the screenshots requested in the assignment.
- `app/data` is the full corpus of 1000 documents used for the final run.

This separation was intentional. It allowed quick debugging and screenshot collection on 100 documents while still satisfying the final assignment requirement of indexing at least 1000 documents.

1.2 Dataset and Document Format

The solution uses a prepared subset of Wikipedia plain text documents. Each document is stored as a UTF-8 text file and follows the required naming convention:

`<doc_id>_<doc_title>.txt`

Examples of valid document names include:

- `3097302_A_Kurt_Weill_Cabaret.txt`
- `18167014_A_British_Picture.txt`

- 34589053_A_Fugitive_from_the_Past.txt

The document id is taken from the prefix before the first underscore. The remaining part of the filename is treated as the document title. The document body is stored as plain text inside the file. The solution does not depend on XML or JSON files.

For grading stability, the repository ships with the prepared local corpora in `app/data` and `app/data_100`. At the same time, the repository also includes a reproducible parquet-based preparation path. The scripts `prepare_corpus_from_parquet.py` and `prepare_corpus_from_parquet.sh` can read a parquet file with Spark, recreate the local UTF-8 text documents, upload them to HDFS, and continue with the standard indexing preparation flow. This means the repository supports both a ready-to-run text corpus and a formal parquet-to-text data preparation path.

1.3 Data Preparation in HDFS with PySpark

The assignment requires a preparation step that creates an HDFS input of the following form:

<doc_id> <doc_title> <doc_text>

This was implemented with two layers of scripts.

The first layer is optional and reproduces the local text corpus from parquet:

- `prepare_corpus_from_parquet.sh`
- `prepare_corpus_from_parquet.py`

`prepare_corpus_from_parquet.sh` performs the following steps:

- Upload the source parquet file to HDFS.
- Run a PySpark job that reads the parquet file.
- Select the `id`, `title`, and `text` columns.
- Create UTF-8 text files with the required <doc_id>_<doc_title>.txt naming format.
- Store the recreated corpus in a local generated folder.
- Continue with the standard HDFS preparation flow.

The second layer prepares the HDFS indexing input from the local text corpus:

- `prepare_data.sh`
- `prepare_data.py`

`prepare_data.sh` performs the orchestration work:

- It receives the local corpus folder and the HDFS source and output paths.
- It deletes old HDFS output if it already exists.
- It uploads the local corpus into HDFS.
- It calls `prepare_data.py` with `spark-submit`.

The upload step was implemented in batches instead of calling `hdfs dfs -put` once per file. This change was important for the final full run on 1000 documents because per-file uploads made the bootstrap much slower.

`prepare_data.py` uses PySpark RDD operations:

- `wholeTextFiles` reads the text files from HDFS.
- The document filename is parsed into `doc_id` and `doc_title`.
- Document text is normalized by collapsing repeated whitespace.
- Invalid or empty records are filtered out.
- The output is coalesced into one partition and stored in HDFS.

As a result, the prepared indexing input is stored as a single `part-00000` file in:

- `/input/data_100` for the 100-document debug run
- `/input/data` for the 1000-document final run

1.4 Text Normalization and Tokenization

The same text processing rules are used both during indexing and during query execution. This is important because the search engine would otherwise tokenize documents and queries differently, which would reduce retrieval quality.

The shared logic is implemented in `text_utils.py`:

- Repeated whitespace is normalized to a single space.
- All tokens are lowercased.
- Tokenization uses a Unicode-aware regular expression:

```
[^\W_]+(?:&#x27;[^\W_]+)*
```

- Apostrophes inside words are preserved.
- Underscores are not treated as token characters.

The system intentionally does not use stemming or stop-word removal. This keeps the baseline solution simple, deterministic, and easy to explain. It also stays close to the assignment requirement, which focuses on distributed indexing and ranking rather than advanced NLP preprocessing.

1.5 Hadoop Streaming Index Construction

The index is built by `create_index.sh`. The script locates the Hadoop Streaming jar and launches four MapReduce jobs. Each job reads the prepared input file from HDFS and produces one search-related artifact.

The final HDFS index layout is:

- `/indexer/documents`
- `/indexer/postings`
- `/indexer/vocabulary`
- `/indexer/corpus_stats`

Each output directory is written with a single reducer so the final artifact is stored as one partition. This design simplifies inspection, loading to Cassandra, and report demonstration.

Job 1: Document Metadata

mapper1.py parses the prepared input and emits:

```
doc_id doc_title doc_len
```

where doc_len is the number of tokens in the document body after tokenization.

reducer1.py keeps one line per document id. The output of this job is used as the document metadata table.

Job 2: Postings

mapper2.py parses each prepared record and computes the term frequency (tf) of every unique term in the document. For each term it emits:

```
term doc_id tf doc_len
```

reducer2.py forwards the grouped output. This job creates the inverted index postings that are later stored in Cassandra.

Job 3: Vocabulary and Document Frequency

mapper3.py emits one line per unique term inside each document:

```
term 1
```

reducer3.py sums these values and produces:

```
term df
```

where df is the document frequency, meaning the number of documents that contain the term.

Job 4: Corpus Statistics

mapper4.py emits corpus-wide aggregate contributions from each document:

- doc_count 1
- total_doc_length doc_len

reducer4.py sums both values and computes:

- doc_count
- total_doc_length
- avg_doc_length

These statistics are required by BM25 at query time.

1.6 Cassandra Schema

The final index is loaded into Cassandra in the keyspace:

```
simple_search_engine
```

The schema contains four tables.

documents

Stores one row per document:

- doc_id
- title
- display_title
- doc_len

display_title is simply the human-readable version of the filename title, with underscores replaced by spaces.

vocabulary

Stores one row per term:

- term
- df

This table is used to retrieve document frequency values during BM25 scoring.

postings

Stores the inverted index:

- term
- doc_id
- tf
- doc_len

The primary key is (term, doc_id), which is a natural access path for query evaluation because the ranker fetches postings by term.

corpus_stats

Stores the small set of global values:

- doc_count
- total_doc_length
- avg_doc_length

This table keeps the BM25 corpus parameters available without recomputing them at query time.

1.7 Loading the Index into Cassandra

The Cassandra loading step is handled by:

- store_index.sh
- app.py

`store_index.sh` validates that all required HDFS index paths exist and then calls `python app.py load /indexer`.

`app.py` performs the following tasks:

- Connects to Cassandra, waiting until the service becomes ready.
- Creates the keyspace and tables if they do not exist.
- Truncates old tables so the load is repeatable.
- Streams each HDFS artifact with `hdfs dfs -cat`.
- Inserts the rows into Cassandra in chunks with prepared statements and concurrent batched execution.

This design keeps the loading logic simple and robust. It also makes the pipeline repeatable without manual cleanup.

1.8 Query Processing and BM25 Ranking

The search engine query path is implemented by:

- `search.sh`
- `query.py`

`search.sh` runs `query.py` on YARN with `spark-submit` and ships the packed Python virtual environment to executors.

`query.py` performs the ranking process in the following order:

- Read the input query from CLI arguments or standard input.
- Tokenize the query with exactly the same rules as the document tokenizer.
- Fetch `doc_count` and `avg_doc_length` from Cassandra.
- Fetch `df` values for all query terms from `vocabulary`.
- Fetch postings for all query terms from `postings`.
- Parallelize those postings as an RDD in Spark.
- Compute BM25 contribution per `(term, document)` pair.
- Sum contributions by document id.
- Return the top 10 documents by decreasing score.

The implemented BM25 form is:

$$\text{score}(d, q) = \text{sum over query terms of } \text{idf}(t) * ((k1 + 1) * \text{tf}(t,d)) / (k1 * ((1 - b) + b * dl/avgdl) + \text{tf}(t,d))$$

with the following hyperparameters:

- $k1 = 1.5$
- $b = 0.75$

The ranker also multiplies the contribution by query term frequency so repeated terms in the query receive more weight.

1.9 Important Design Choices

Several design decisions were made deliberately:

- The repository uses four simple MapReduce jobs instead of one complex job. This makes the index artifacts easier to inspect and explain.
- The full system is repeatable. Running indexing again overwrites HDFS outputs and truncates Cassandra tables.
- Query-time ranking reads only the postings needed for the query terms. The system does not scan all documents during search.
- Titles are stored for output only. Ranking is based on document text, not title text.
- A fixed debug subset of 100 documents was kept in the repository because the assignment explicitly requires screenshots for a 100-document indexing run.

The main limitation of the system is that query processing still collects postings for the query terms before parallel scoring. This is acceptable for a simple academic baseline and keeps the implementation easy to understand.

2. Demonstration

2.1 Repository Startup

The repository can be started with Docker Compose.

For the final full run:

```
docker compose up
```

For the debug run used in the screenshots:

```
docker compose down --remove-orphans  
DATASET_MODE=debug docker compose up
```

In PowerShell, the debug run can be started as:

```
docker compose down --remove-orphans  
$env:DATASET_MODE='debug'  
docker compose up
```

The main entrypoint is `app/app.sh`. It starts the Hadoop services, reuses or creates the Python virtual environment, prepares the chosen corpus, runs indexing, loads Cassandra, and executes sample queries.

If a parquet file is available, the same repository can reproduce the local corpus before indexing. For example, if `app/a.parquet` is present, the full run can be started with:

```
$env:PARQUET_INPUT_PATH='/app/a.parquet'  
docker compose up
```

and the debug run on 100 recreated documents can be started with:

```
$env:DATASET_MODE='debug'  
$env:PARQUET_INPUT_PATH='/app/a.parquet'  
docker compose up
```


2.2 Debug Demonstration on 100 Documents

The debug scenario exists for two reasons:

- It satisfies the assignment requirement for screenshots of a successful 100-document indexing run.
- It provides a smaller and faster environment for validation before scaling to 1000 documents.

Figure S01. Successful debug bootstrap

This screenshot shows the successful end-to-end bootstrap for the debug mode. The environment starts correctly, the data is prepared in HDFS, the index is built and loaded, and sample search queries are executed.



```
1 # big-data-assignment2
2
3 ## Current repository state
4
5 This repository now includes a working end-to-end baseline solution for the assignment:
6
7 1. Hadoop and Cassandra services are started with Docker Compose.
8 2. A local text corpus is uploaded to HDFS.
9 3. PySpark prepares /input/data or /input/data_100 as one partition.
10 4. Hadoop Streaming MapReduce jobs build:
11   - document metadata
12   - postings
13   - vocabulary with document frequencies
14   - corpus statistics for BM25
15 The index is loaded into Cassandra tables.
16 6. search.sh runs a PySpark BM25 raster on YARN and prints top results.
17
18 PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
19
20 cluster-master 26/04/05 09:45:35 INFO BlockManager: Using org.apache.spark.storage.BlockManager$DefaultTopology for block replication policy
21 cluster-master 26/04/05 09:45:35 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, cluster-master, 42689, None)
22 cluster-master 26/04/05 09:45:35 INFO BlockManagerMasterEndpoint: Registering block manager cluster-master-42689 with 365.3 MB RAM, BlockManagerId(driver, cluster-master, 42689, None)
23 cluster-master 26/04/05 09:45:35 INFO BlockManager: Registered BlockManager: BlockManagerId(driver, cluster-master, 42689, None)
24 cluster-master 26/04/05 09:45:35 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, cluster-master, 42689, None)
25 cluster-master 26/04/05 09:45:35 INFO HadoopClientSchedulerEndpoint: Add Hadoop Filter: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter, Map(PROXY_HDSTS -> cluster-master, PROXY_H1_BUGS -> http://cluster-master:8080/proxy/application_177530048056_0007/, /proxy/application_177530048056_0007)
26 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /jobs: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
27 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /jobs/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
28 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /stages: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
29 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /stages/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
30 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /stages/pool: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
31 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /storage: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
32 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /storage/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
33 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /storage/pool: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
34 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /environment: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
35 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /environment/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
36 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /executors: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
37 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /executors/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
38 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /executors/threadDump: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
39 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /static: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
40 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
41 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /fail: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
42 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /stages/stopKill: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
43 cluster-master 26/04/05 09:45:35 INFO ServerInfo: Adding filter to /metrics/job: org.apache.hadoop.yarn.server.webproxy.authfilter.AuthFilter
44 cluster-master 26/04/05 09:45:35 INFO HadoopClientSchedulerEndpoint: ApplicationMaster registered as http://cluster-master:8080/proxy/application_177530048056_0007/ with ID 1, ResourceProfileId 0
45 cluster-master 26/04/05 09:45:35 INFO HadoopClientSchedulerEndpoint: Registered executor http://cluster-master:8080/proxy/application_177530048056_0007/ with ID 2, ResourceProfileId 0
46 cluster-master 26/04/05 09:45:35 INFO HadoopClientSchedulerEndpoint: Scheduler is ready for scheduling beginning after reaching minRegisteredMemoryRatio: 0.0
47 cluster-master 26/04/05 09:45:35 INFO BlockManagerMasterEndpoint: Registering block manager cluster-master-145625 with 365.3 MB RAM, BlockManagerId(cluster-master-1, 41675, None)
48 cluster-master Top 10 results for query: death
49 cluster-master Rank doc_id title score
50 cluster-master 1 1822877 A Death in the Family (comic) 3.64212
51 cluster-master 2 1086660 A Day in the Death of Drury 9 3.23612
52 cluster-master 3 1083136 A Decade in the Grove 3.08813
53 cluster-master 4 1223771 A Day in the Death 3.08715
54 cluster-master 5 1180274 A Cold Night's Death 2.85741
55 cluster-master 6 1025482 A Flat Run 2.64845
56 cluster-master 7 145793 A Double Life (1947 film) 2.180558
57 cluster-master 8 1348184 A Gamble with Death 2.73818
58 cluster-master 9 3957006 A Gun Called Tomlin 2.64030
59 cluster-master 10 12547743 A Darkness More Than Night 2.48801
60 cluster-master Bootstrap complete.
61 cluster-master Prepared index is ready in HDFS: /input/data_100
62 cluster-master Index base is ready in Cassandra keyspace: simple_search_engine
```

Figure S01. Successful bootstrap of the debug run on 100 documents.

Figure S02. Successful indexing output for 100 documents

This screenshot is the key indexing proof for the report. It shows the HDFS counts for the raw corpus, the prepared input, and the four index artifacts. It also shows the corpus statistics output including doc_count = 100.




Figure S02. HDFS counts and corpus statistics for the 100-document debug index.

Figure S03. Debug search for history

The history query demonstrates that the engine is able to retrieve and rank documents using BM25 instead of just filtering exact matches. The output contains document ids, human-readable titles, and numeric scores.





Figure S03. Top 10 search results for the query history in debug mode.

Figure S04. Debug search for dogs

The dogs query demonstrates another query type with a concrete thematic keyword. The returned documents are ranked by term frequency, document frequency, and document length normalization. Since the engine

ranks by document body and not by title, the title alone does not always reveal why the document scored highly. This is expected behavior and not an error.




```
1 # big-data-assignment2
2
3 ## Current repository state
4
5 This repository now includes a working end-to-end baseline solution for the assignment:
6
7 1. Hadoop and Cassandra services are started with Docker Compose.
8 2. A local text corpus is uploaded to HDFS.
9 3. PySpark prepares '/input/data' or '/input/data_100' as one partition.
10 4. Hadoop Streaming MapReduce jobs build:
11   - document metadata
12   - postings
13   - vocabulary with document frequencies
14   - corpus statistics for BM25
15 5. The index is loaded into Cassandra tables.
16 6. 'search.sh' runs a PySpark BM25 ranker on YARN and prints top results.
17
18 PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
19
20 PS C:\Users\ve11\OneDrive\Public> cd ~/big-data-assignment2; docker exec cluster-master bash -li 'cd /app && bash search.sh "dogs"'
21 start time: 17753832347
22 final status: OK[7/10]
23 tracking URL: http://cluster-master:8088/proxy/application_1775382846565_0000/
24 user: root
25 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Application application_1775382846565_0000 has started running.
26 26/04/05 18:08:46 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 42597.
27 26/04/05 18:08:46 INFO NettyBlockTransferService: Server created on cluster-master/42597
28 26/04/05 18:08:46 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
29 26/04/05 18:08:46 INFO BlockManagerMasterEndpoint: Registering block manager cluster-master/42597 with 366.3 MiB RAM, BlockManagerId(driver, cluster-master, 42597, Now)
30 26/04/05 18:08:46 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, cluster-master, 42597, Now)
31 26/04/05 18:08:46 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, cluster-master, 42597, Now)
32 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
33 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/jobs: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
34 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/jobs/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
35 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
36 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
37 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/stage: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
38 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/stage/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
39 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/pool: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
40 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/pool/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
41 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
42 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
43 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/rdd: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
44 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/rdd/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
45 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /newcommand/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
46 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /newcommand: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
47 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
48 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
49 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors/heapdump: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
50 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stats: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
51 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stats/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
52 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/tagkill: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
53 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /metrics/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
54 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Registered executor NettyBlockTransferService (spark-client//executor) (172.18.0.3:42000) with ID 1, ResourceProfileId 0
55 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Registered executor NettyBlockTransferService (spark-client//executor) (172.18.0.3:42000) with ID 2, ResourceProfileId 0
56 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minResourcesThreshold: 6.0
57 26/04/05 18:08:51 INFO BlockManagerMasterEndpoint: Registered block manager cluster-slave1-14801 with 366.3 MiB RAM, BlockManagerId(2, cluster-slave-1, 44001, Now)
58 Top 10 results for query: dogs
59 rank doc_id title score
60 1 10658697 A Dangerous Path 8.278025
61 2 143450 A Day In the Life 1.84668
62
63 PS C:\Users\ve11\OneDrive\Public> cd ~/big-data-assignment2;
```

Figure S04. Top 10 search results for the query dogs in debug mode.

Figure S05. Debug search for film

The film query demonstrates a different category of term. In this case, the results are especially useful because many Wikipedia articles in the corpus are about movies, TV series, actors, or entertainment-related topics. This makes the retrieval behavior easier to understand visually from the titles alone.



```
1 # big-data-assignment2
2
3 ## Current repository state
4
5 This repository now includes a working end-to-end baseline solution for the assignment:
6
7 1. Hadoop and Cassandra services are started with Docker Compose.
8 2. A local text corpus is uploaded to HDFS.
9 3. PySpark prepares '/input/data' or '/input/data_100' as one partition.
10 4. Hadoop Streaming MapReduce jobs build:
11   - document metadata
12   - postings
13   - vocabulary with document frequencies
14   - corpus statistics for BM25
15 5. The index is loaded into Cassandra tables.
16 6. 'search.sh' runs a PySpark BM25 ranker on YARN and prints top results.
17
18 PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
19
20 PS C:\Users\ve11\OneDrive\Public> cd ~/big-data-assignment2; docker exec cluster-master bash -li 'cd /app && bash search.sh "film"'
21 start time: 17753832347
22 final status: OK[7/10]
23 tracking URL: http://cluster-master:8088/proxy/application_1775382846565_0000/
24 user: root
25 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Application application_1775382846565_0000 has started running.
26 26/04/05 18:08:46 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 44835.
27 26/04/05 18:08:46 INFO NettyBlockTransferService: Server created on cluster-master/44835
28 26/04/05 18:08:46 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
29 26/04/05 18:08:46 INFO BlockManagerMasterEndpoint: Registering block manager cluster-master/44835 with 366.3 MiB RAM, BlockManagerId(driver, cluster-master, 44835, Now)
30 26/04/05 18:08:46 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, cluster-master, 44835, Now)
31 26/04/05 18:08:46 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, cluster-master, 44835, Now)
32 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
33 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/jobs: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
34 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /jobs/jobs/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
35 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
36 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
37 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/stage: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
38 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/stage/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
39 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/pool: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
40 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/pool/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
41 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
42 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
43 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/rdd: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
44 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/rdd/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
45 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /newcommand/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
46 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /newcommand: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
47 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
48 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
49 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /executors/heapdump: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
50 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stats: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
51 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stats/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
52 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /stages/tagkill: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
53 26/04/05 18:08:46 INFO ServerInfo: Adding filter to /metrics/json: org.apache.hadoop.yarn.server.webproxy.filter.AclFilter
54 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Registered executor NettyBlockTransferService (spark-client//executor) (172.18.0.3:41356) with ID 1, ResourceProfileId 0
55 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: Registered executor NettyBlockTransferService (spark-client//executor) (172.18.0.3:41356) with ID 2, ResourceProfileId 0
56 26/04/05 18:08:46 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minResourcesThreshold: 6.0
57 26/04/05 18:08:47 INFO BlockManagerMasterEndpoint: Registered block manager cluster-slave1-14347 with 366.3 MiB RAM, BlockManagerId(1, cluster-slave-1, 43437, Now)
58 26/04/05 18:08:47 INFO BlockManagerMasterEndpoint: Registered block manager cluster-slave1-14373 with 366.3 MiB RAM, BlockManagerId(2, cluster-slave-1, 45737, Now)
59 26/04/05 18:08:47 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minResourcesThreshold: 6.0
60 Top 10 results for query: film
61 rank doc_id title score
62 1 10534786 A Black and White World 2.048809
63 2 1368587 A Fascinal of Time 1.995442
64 3 1500666 A Final Reducing 1.932846
65 4 146993 A Double Life (1947 Film) 1.913069
66 5 1517069 A Gentle Peace in the Village 1.908185
67 6 1022157 A Balinese France Source 1.901703
68 7 1121173 A Gentleman of Paris (1927 film) 1.899775
69 8 1466166 A Day at the Beach 1.899278
70 9 1121461 A Festival of Frigate 1.888657
71 10 1040480 A Day In the Death of Danny B 1.837902
72
73 PS C:\Users\ve11\OneDrive\Public> cd ~/big-data-assignment2;
```

Figure S05. Top 10 search results for the query film in debug mode.

2.3 Final Demonstration on 1000 Documents

After validating the system on the 100-document debug subset, the exact same pipeline was executed on the full local corpus of 1000 documents.

The default `docker compose up` command now runs this final mode automatically. This is important because the evaluator is expected to clone the repository and start it with Docker Compose.

Figure S06. Successful full bootstrap

This screenshot shows the successful final bootstrap on the full corpus. It confirms that the repository is not limited to the debug subset.

```
File Edit Selection View Go Run Terminal Help
@ README.md X
1 # big-data-assignment2
2 # Current repository state
3
4 This repository now includes a working end-to-end baseline solution for the assignment:
5
6 1. Hadoop and Cassandra services are started with Docker Compose.
7 2. A local test corpus is uploaded to HDFS.
8 3. PySpark prepares /input/data or /input/data_100 as one partition.
9 4. Hadoop Streaming MapReduce job builds:
10   - document metadata
11   - postings
12   - vocabulary with document frequencies
13   - corpus statistics for BM25
14
15 The index is loaded into Cassandra tables.
16
17 6. 'search.sh' runs a PySpark BM25 ranker on YARN and prints top results.
18
19 By default, 'docker compose up' will:
20
21 1. Start the Hadoop master, Hadoop worker, and Cassandra containers.
22 2. Create or reuse the Python virtual environment inside the master container.
23 3. Upload the full corpus from 'app/data' to HDFS as '/data'.
24 4. Build the prepared input dataset in HDFS as '/input/data'.
25 5. Build the index in HDFS and store it in Cassandra.
26 6. Run a few sample search queries.
27
28
29 PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
30
31 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /jobs/join: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
32 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /jobs/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
33 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /jobs/job/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
34 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /stages/stage: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
35 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /stages/stage/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
36 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /stages/pool: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
37 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /storage/pool: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
38 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /storage/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
39 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /storage/pool: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
40 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /storage/pool/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
41 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /environment: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
42 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /environment/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
43 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /executors/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
44 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /executors/threadname: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
45 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /executors/threadname/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
46 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /static: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
47 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
48 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /api: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
49 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /info/info: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
50 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /stages/stage/kill: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
51 cluster-master 26/04/05 18:28:17 INFO ServerInfo: Adding filter to /metrics/job: org.apache.hadoop.yarn.server.webproxy.filter.AadFilter
52 cluster-master 26/04/05 18:28:17 INFO YarncheddierBackendYarnDriverEndpoint: Registered executor NettyRpcEndpointRef(spark-client://Executor) (172.18.0.2:3804) with ID 1, ResourceProfileId 0
53 cluster-master 26/04/05 18:28:18 INFO BlockManagerMasterEndpoint: Registering block manager cluster-slave-144802 with 366.3 MB RAM, BlockManagerId: (cluster-slave-1, 44802, None)
54 cluster-master 26/04/05 18:28:18 INFO YarncheddierBackendYarnDriverEndpoint: Registered executor NettyRpcEndpointRef(spark-client://Executor) (172.18.0.2:3808) with ID 2, ResourceProfileId 0
55 cluster-master 26/04/05 18:28:18 INFO BlockManagerMasterEndpoint: Registering block manager cluster-slave-144802 with 366.3 MB RAM, BlockManagerId: (cluster-slave-1, 44802, None)
56 cluster-master 26/04/05 18:28:18 INFO YarncheddierBackendYarnDriverEndpoint: SchedulerBackend is ready for scheduling beginning after reached minResourceDuration: 0.0
57 cluster-master Top 10 results for query: death
58 cluster-master Rank Doc ID Title Score
59 cluster-master 1 15864758 A Letter from Death Row 4.124245
60 cluster-master 2 4081128 A Bad Called Death 4.087085
61 cluster-master 3 1822677 A Death in the Family (comic) 4.059913
62 cluster-master 4 808157 A Bad A Q (album) 3.753128
63 cluster-master 5 7136475 A Death in Bed No. 12 3.742843
64 cluster-master 6 5723123 A Death of Innocence 3.055094
65 cluster-master 7 1888968 A Day in the Death of Dory B 3.002487
66 cluster-master 8 10942176 A Day in Summer 3.052111
67 cluster-master 9 1381136 A Decade in the Grave 3.148810
68 cluster-master 10 2054208 A Blaze in the Northern Sky 3.130701
69 cluster-master
70 cluster-master Bootstrap complete.
71 cluster-master Prepared input is ready in HDFS: /input/data
72 cluster-master Index base is ready in Cassandra keyspaces: simple_search_engine
```

Figure S06. Successful bootstrap of the full 1000-document run.

Figure S07. HDFS counts and full corpus statistics

This screenshot shows that the final HDFS structures were created correctly for the full run. It also confirms that the corpus statistics were computed for all 1000 documents.




Figure S07. HDFS counts and corpus statistics for the 1000-document full index.

From the verified final run, the following values were produced:

- doc_count = 1000
- avg_doc_length = 573.655
- vocabulary size = 42773
- postings count = 252622

These numbers show that the final full index is substantially larger than the debug index and that the pipeline scales correctly from the 100-document subset to the 1000-document corpus.

Figure S08. Full search for history of music

The final query history of music demonstrates a multi-term search in the full corpus. This is a better test than a single-term query because it verifies that the ranker sums BM25 contributions across several query terms.


```
docker compose up
```

Manual indexing from inside the master container

```
docker exec -it cluster-master bash
cd /app
bash index.sh /input/data
```

Manual search from inside the master container

```
bash search.sh "history of music"
```

4. Conclusion

The final repository implements the complete baseline required by the assignment:

- Plain text documents are stored with the required filename format.
- Documents are uploaded to HDFS and converted into a prepared one-partition indexing input with PySpark.
- Hadoop Streaming MapReduce builds the document metadata, postings, vocabulary, and corpus statistics.
- Cassandra stores the final index structures.
- PySpark on YARN computes BM25 and returns the top 10 ranked documents for a query.

The debug run on 100 documents and the final run on 1000 documents both completed successfully. The screenshots included in the PDF report demonstrate successful indexing and successful query execution in both modes. The solution therefore satisfies the functional requirements of the assignment while keeping the implementation transparent and easy to validate.